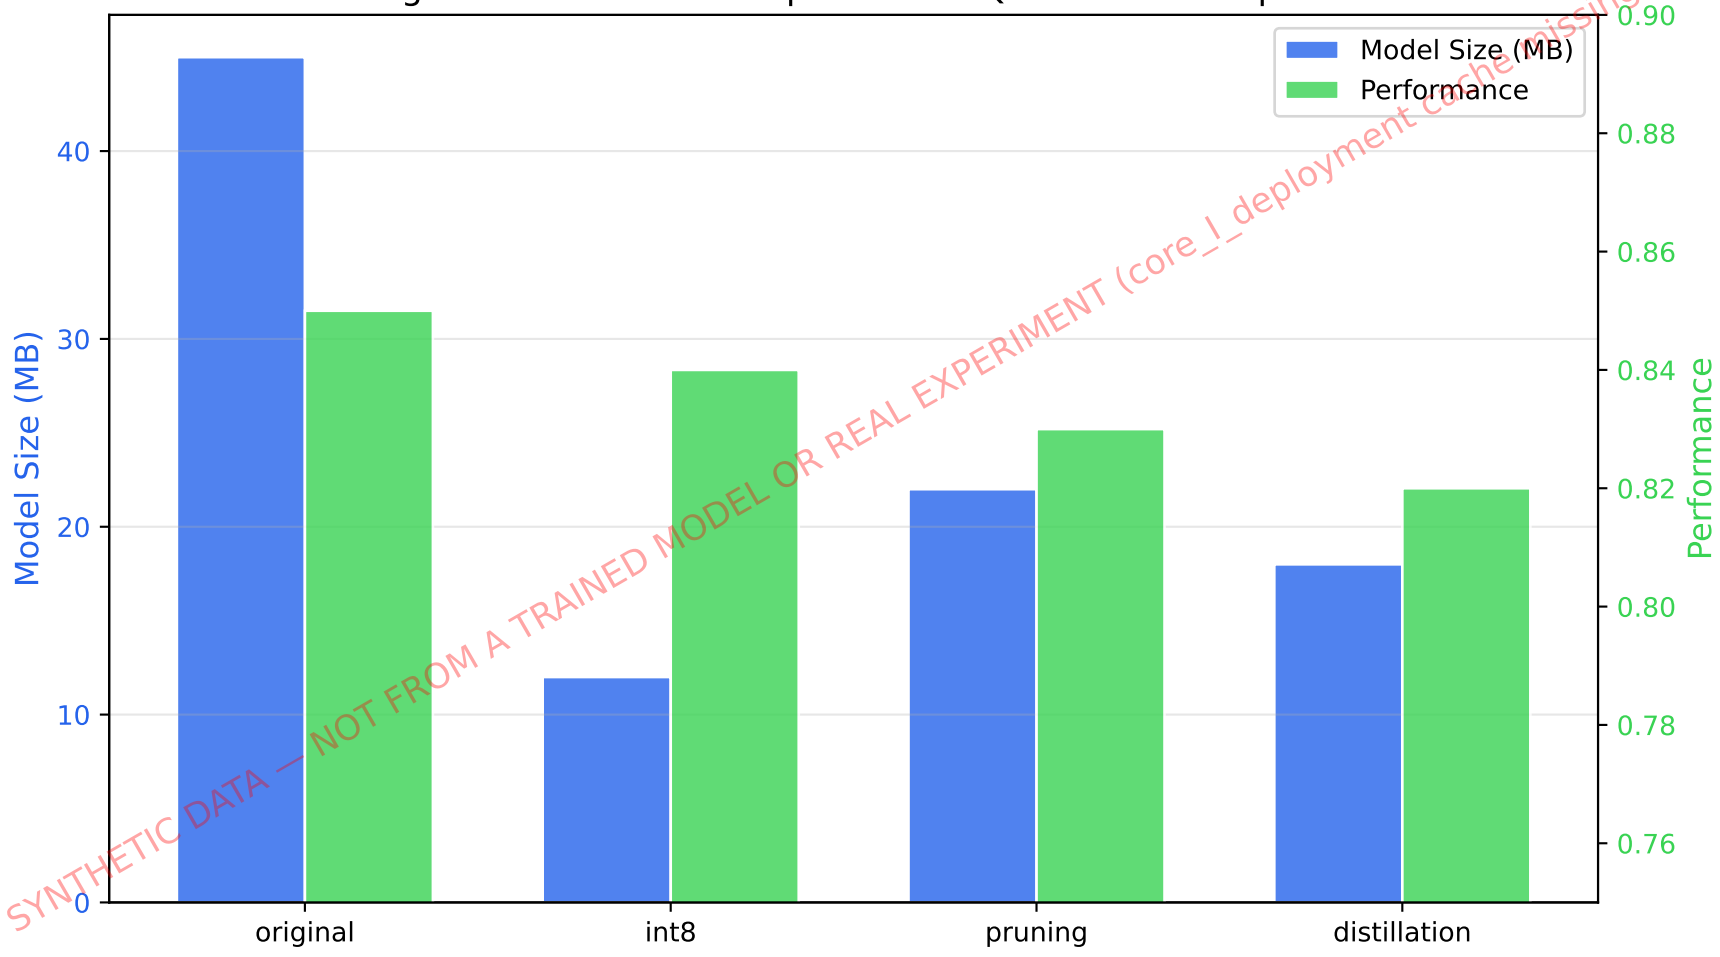
Figure 53 — Model Compression & Quantization Impact
Model Size (MB)
Performance
Model Size (MB)
Performance
0
10
20
30
40
0.76
0.78
0.80
0.82
0.84
0.86
0.88
0.90
original
int8
pruning
distillation
SYNTHETIC DATA — NOT FROM A TRAINED MODEL OR REAL EXPERIMENT (core_l_deployment cache missing)